

Knowledge Assistant Retrieval Test

Test document for text extraction, simple table extraction, chunking, retrieval, and grounded answers.

1. Product Overview

Atlas Support Hub is an internal knowledge assistant for the customer operations team. It helps employees find approved information about support procedures, service levels, and escalation rules from uploaded company documents.

The assistant must answer only from the approved documents. If the requested information is not present, it must say: I could not find this information in the uploaded documents.

2. Service Configuration

Area	Configured Value	Reason
Support Hours	Monday to Friday, 09:00 to 18:00	Standard business-hour coverage
Priority One SLA	Initial response within 30 minutes	Critical incidents require immediate attention
Priority Two SLA	Initial response within 4 hours	Important issues are handled during the same business day
Priority Three SLA	Initial response within 2 business days	General requests can follow the normal support queue

3. Escalation Rules

Priority	Escalation Owner	Escalation Time	Required Action
Priority One	Incident Manager	After 30 minutes without progress	Open an incident bridge and notify the Operations Lead
Priority Two	Support Team Lead	After 4 business hours	Review the ticket and assign a technical owner
Priority Three	Support Coordinator	After 2 business days	Recheck the queue and update the requester

4. Document Processing Policy

Uploaded documents are processed in the following order: text extraction, whitespace cleanup, chunk creation, embedding generation, vector storage, and retrieval during chat. Each answer should include the document name and page number when a source is available.

Only PDF, TXT, and Markdown files are supported for the current MVP. Spreadsheet uploads are not supported by this policy.

5. Approved Technology Choices

Area	Preferred Option	Reason
Backend	Python and FastAPI	Simple API development and a strong AI ecosystem
Vector Database	Qdrant	Open-source vector search with metadata support
Embeddings	nomic-embed-text	Local embedding generation without a paid API
Language Model	Ollama	Local answer generation for the MVP

Grounding rule: The assistant must not invent a policy, SLA, owner, or technology choice that is not written in this document.